BUILT · WORKING

Reading the paper `ocr/`

A research paper is a PDF — a picture of a document, not data. This stage turns it into text a program can actually work with.

WHAT GOES IN, WHAT COMES OUT

INPUT	WHAT HAPPENS	OUTPUT
A paper PDF e.g. “Network In Network”, 10 pages	**Each page → an image → Claude reads it** one call per page, ~13 seconds each	**One Markdown file** plain text, tables, figure descriptions

THE KEY POINTS

- **It looks at the page, it doesn't parse the file.** Every page is rendered as an image and read by Claude — the same way a person reads a printed paper.
- **Layout is handled deliberately.** Two-column papers are read down each column rather than straight across, and page numbers, headers and the arXiv stamp are dropped instead of landing mid-sentence.
- **Tables are copied cell by cell.** They hold the results the paper is claiming, so a single wrong digit here would quietly corrupt every later stage.
- **Figures come out as sentences** — which is where the whole approach earns its keep, for the reason shown below.
- **About two minutes per paper.** Roughly 13 seconds a page, one call each.

WHY THAT FOURTH POINT MATTERS

FROM THE OUTPUT — A FIGURE, AS TEXT

```
[Figure 2: The overall structure of Network In
Network. The NINs include the stacking of
three mlpconv layers and one global average
pooling layer…]
```

That sentence is the paper's entire architecture. It exists *only* in the figure — no table or paragraph states it. A text-scraping tool would have missed it, and the stage that later writes the training code would have had nothing to build from.

WHERE IT STANDS

6 / 8
papers converted

320 KB
of clean text

2
blocked — below

The one real problem. Two papers refuse to convert. Both fail on page 2, every time, because Claude's safety filter flags the page image — even though the pages are ordinary academic text. We checked: it fails identically at a different image size and with a one-line prompt, so it isn't something we're doing wrong. Those two are stuck outside the pipeline for now.