

ReproBot: A Multi-Agent System for Automated Scientific Paper Replication

First Progress Report

Aleyna Kütük
aleynakutuk10@gmail.com

Engincan Varan
engincanvaran@gmail.com

Gül Korkut
serifegulkorkut@gmail.com

Göktuğ Mert Özdoğan
goekmeroz@gmail.com

inzva AI Projects #10
17.07.2026

Abstract

Reproducing the empirical results of a machine learning paper remains a largely manual, labor-intensive task, with the large majority of published work shipping no usable implementation at all. **ReproBot** is a proposed multi-agent pipeline that reads an ML paper in PDF form, extracts its method and claims, generates and executes a training script, and produces a structured report comparing reproduced metrics against the paper’s stated numbers – retrying the implementation when a Critic agent detects a mismatch. This document reports on the project’s first phase: a literature review of nine closely related agentic-ML and paper-replication systems, a detailed implementation-ready project plan (architecture, agent specifications, evaluation rubric, timeline, risk register, and cost budget), and an early prototype of the Reader agent that drafts method summaries, claims, and hyperparameters from paper PDFs using the Anthropic API together with `pdfplumber`-based text extraction – so far tried only on a small sample of CIFAR-10 papers. We summarize what has been completed, what remains open, and the concrete next steps for maturing the Reader and beginning the Coder and Runner components.

1 Introduction and Motivation

Reproducing the empirical results of a machine learning paper is, in principle, a mechanical exercise: read the method, implement it, run it, and compare the numbers. In practice it is one of the most labor-intensive tasks in ML research. A reader must reconstruct architectural details that authors often under-specify, infer hyperparameters and data-processing conventions treated as “common knowledge” within a subfield, and debug an implementation against a single, sparse signal – whether the final reported metric matches. The problem is compounded by how rarely authors release usable code:

across ICLR, ICML, and NeurIPS 2024, only 16.7–21.2% of accepted papers shipped a working implementation [3], leaving the large majority of the field’s claims effectively unverifiable without a from-scratch reimplementation.

Large language model (LLM) agents have, over the last two years, been applied to nearly every stage of this problem individually: benchmarked on general ML experimentation [1], tasked with generating entirely new research ideas and writing them up as papers [7, 9], and used as coding assistants that translate a paper’s method section into a code repository [3]. A rigorous benchmark now exists for measuring paper-replication capability directly, complete with author-validated grading rubrics [2]. Yet across this body of work, a specific combination of capabilities – parsing a paper’s claims, executing a faithful implementation, and explicitly verifying that the reproduced numbers match what was published, with an iterative repair loop driven by that comparison – has not been assembled into a single system (Section 2).

ReproBot targets exactly this combination. It is structured as a four-agent pipeline – Reader, Coder, Runner, and Critic – coordinated by a central Orchestrator over a shared memory state (Section 3). The Reader ingests a paper PDF and extracts a structured method summary, claims, datasets, and hyperparameters; the Coder translates this into a self-contained HuggingFace **Trainer**-based training script; the Runner executes the script in an isolated sandbox and captures metrics and error traces; and the Critic compares the reproduced metrics against the paper’s stated numbers, producing a pass/retry/fail verdict and, on retry, targeted feedback for the Coder. The pipeline is deliberately scoped at launch to a benchmark of 20 image-classification papers, trading the topic breadth of general-purpose benchmarks for a tighter, more auditable definition of “replication success” within a single, well-understood domain.

This report covers the project’s first phase of work:

literature review, detailed planning, and an early Reader agent prototype. Section 2 reviews the related work that shaped ReproBot’s design. Section 3 describes the proposed architecture. Section 4 describes the evaluation methodology. Section 5 lays out the project timeline. Section 6 is the core contribution of this report: a concrete account of what has been completed to date and what remains open. Section 7 describes next steps, and Section 8 concludes.

2 Related Work

2.1 Foundations: Benchmarking Agentic ML Experimentation

The starting point for this line of work is **MLAgentBench** [1], the first benchmark to measure whether LLM agents can conduct open-ended ML experimentation rather than merely propose hyperparameters. It poses 13 executable tasks – spanning canonical benchmarks (CIFAR-10, IMDB sentiment), classic Kaggle competitions, and research-recency-controlled challenges – and scores an agent’s final submission against a starter-code baseline. Its agent architecture, a ReAct-style loop augmented with explicit Reflection, Research-Plan-and-Status, and Fact-Check fields, targets two failure modes the authors identify as central to agentic ML work: poor long-horizon planning and hallucinated progress claims. Even so, the best-performing configuration (Claude 3 Opus) reaches only a 37.5% average success rate across the 13 tasks.

MLAgentBench does not ask an agent to read a paper at all, let alone reproduce one; its contribution to ReproBot is architectural. The finding that unstructured agent loops hallucinate and lose the plot over long horizons directly motivates separating concerns across a Reader, Coder, Runner, and Critic under a single Orchestrator, rather than asking one agent to do everything.

2.2 The Detour Through Autonomous Research Generation

Rather than pursuing replication, the field’s next wave of systems pointed LLM agents at generating *new* research. **The AI Scientist** [7] was the first to close a full loop – idea generation, iterative experimentation, manuscript writing, and automated peer review – end to end, at a reported cost of roughly \$10–15 per paper, but is candid about failure modes directly relevant to any autonomous verification system: models fabricated ablation tables and mischaracterized a regressed metric as an improvement. **The AI Scientist-v2** [8] replaced linear iteration with a stage-gated tree search and added a vision-language model (VLM) critic over generated figures; one of its fully autonomous submissions was

accepted to an ICLR 2025 workshop, though the authors note the other two were rejected.

MLR-Copilot [9] retrieves prototype code and candidate models/datasets directly from HuggingFace rather than generating everything from a blank context, and its cleanest empirical result is directly relevant to ReproBot: a single-shot code-generation baseline succeeds on **0% of trials** across all five evaluated tasks, while its full iterative debugging loop reaches up to 50% success with the strongest backbone tested – quantified evidence that iterative repair, not just code generation, is what makes autonomous implementation viable at all. **Agent Laboratory** [6] rounds out this wave with a three-phase pipeline and a configurable human-in-the-loop mode; its own automated reviewer scores generated papers at 6.1/10 on average versus 3.8/10 from real human reviewers – a 2.3-point overestimation that argues against trusting any pipeline’s self-assessment of its own output without an independent, objective check.

This wave demonstrates that LLM agents can execute the entire research lifecycle end to end, but leaves the replication problem untouched by construction – these systems have no ground-truth number to reproduce, since they invent the research question themselves. They remain complementary to ReproBot, and several of their components – Agent Laboratory’s code-mutation loop with pooled scoring, MLR-Copilot’s retrieve-then-integrate pattern, and The AI Scientist-v2’s stage-gated tree search – are directly transferable design patterns for ReproBot’s own Coder and Critic loop.

2.3 The Replication Wave: Benchmarks and Systems

A third, more recent wave of work returns to the narrower problem ReproBot targets directly: given an existing paper, can an agent reproduce it. **PaperBench** [2] is the field’s gold-standard evaluation instrument. It gives agents 20 ICML 2024 Spotlight/Oral papers and requires a from-scratch code repository, graded against a hierarchical rubric of 8,316 individually gradable leaf-node outcomes split into Code Development, Execution, and Result Match categories. A human baseline shows PhD-level researchers start slower than the best agent configuration but overtake it after roughly 24 hours of sustained work, suggesting current agents plateau on long-horizon tasks rather than lacking raw capability.

PaperCoder [3] decomposes paper-to-code translation into Planning, Analysis, and Coding stages, generating files sequentially in a dependency-aware order. On PaperBench’s lightweight Code-Dev variant it reaches a 45.1% replication score – roughly double the best flat-ReAct baseline – but it **never executes** the code it produces; its own coverage analysis identifies data processing as the weakest-implemented component (56% coverage versus 79–

92% elsewhere), precisely because papers are least explicit about data handling and no execution step exists to surface the gap.

AutoP2C [4] closes the execution gap PaperCoder leaves open by treating paper-to-code as an explicitly multimodal task – parsing prose, equations, tables, and, via a VLM, architecture diagrams – and iteratively debugging the generated code (execute → localize error → correct → repeat) until it runs. It reports the strongest executability numbers surveyed (100% of 8 benchmark papers produce runnable code, versus 12.5% for single-pass o1/DeepSeek-R1 baselines) and an ablation showing that removing architecture-diagram parsing alone costs 21.9 points of absolute reproduced performance. Its one quantitative ground-truth comparison – average “relative performance” of 99.5% – is, however, an aggregate figure with **no stated tolerance or pass/fail criterion**.

AutoReproduce [5] introduces the *paper lineage* algorithm: mining a paper’s citation graph for the papers it most heavily builds on and using their text and linked code as grounding context during generation, paired with a sampling-based unit-testing strategy that dry-runs a training script on a mini-batch before committing to a full run. This dry-run technique alone was the single largest driver of their execution-rate improvement (roughly 2–23% for non-lineage baselines up to 77–95% with lineage). Yet AutoReproduce’s headline metric is execution rate and performance gap against a *reference implementation*, not an explicit comparison against the paper’s own claimed numbers.

2.4 The Gap ReproBot Fills

Read together, the replication wave traces a clean division of labor with one piece missing. PaperCoder proves a paper’s method can become structurally coherent code, but the code is never run. AutoP2C closes the execution gap and adds multimodal parsing, coming closest to ReproBot’s own pipeline shape – but its debugging loop terminates on “the code runs and looks architecturally aligned,” with no stated tolerance on its one ground-truth comparison.

AutoReproduce runs the code and improves on PaperCoder’s benchmark scores via citation-lineage mining, but its success criterion stops at executability and reference-code similarity, not a direct comparison to the paper’s stated claims. **None of the three closes the loop from reproduced metric to an explicit pass/fail verdict against the published number to a targeted retry.** Table 1 makes this concrete across six capabilities: ReproBot is the only system that covers all six.

3 Proposed System Architecture

ReproBot is structured as a four-agent pipeline coordinated by a central **Orchestrator** that owns a single, versioned shared-memory object per paper (Figure 1).

- **Reader.** Ingests a PDF using `pdfplumber` (text/tables) and a vision-language model (figure/diagram parsing) to extract a method summary, claims (metric, dataset, reported value, source location), hyperparameters, and a dedicated data-processing pass – data-processing detail is the single weakest-covered aspect of paper-to-code extraction across every surveyed system (56% coverage in PaperCoder vs. 79–92% elsewhere [3]), so it is treated as a first-class output field rather than folded into the general method summary.
- **Coder.** Translates the Reader’s output into a self-contained HuggingFace **Trainer**-based training script; on retry, applies the Critic’s targeted feedback rather than regenerating from scratch, following MLR-Copilot’s retrieve-before-generate pattern for models and datasets [9].
- **Runner.** Executes the script inside a Docker sandbox with no host access, cheaply validating it first with AutoReproduce’s sampling-based dry run (mini-batch shape check + early-exit full-pipeline pass) before committing to a full training run [5] – this technique alone drove that system’s largest execution-rate gains.
- **Critic.** Compares the Runner’s reproduced metrics against the Reader’s extracted claims and produces a structured **pass/retry/fail** verdict with an explicit numeric tolerance, plus targeted natural-language feedback on retry – the tolerance gate that Section 2.3 identifies as missing from every surveyed replication system.
- **Orchestrator.** Owns the shared-memory object, enforces the retry budget and confidence threshold, and checks the append-only history log before each retry to avoid burning the budget on a plateaued fix.

Model selection standardizes on the Claude family: Claude Sonnet 5 as the default backbone for the Reader, Coder, Critic, and Report Generator, with Claude Opus 4.8 reserved for low-confidence extractions, stuck retries (2+ failed attempts), and borderline numeric judgment calls, and Claude Haiku 4.5 for cheap, high-volume classification (Runner outcome triage, Orchestrator edge-case routing). This choice is grounded directly in PaperBench’s own backbone comparison, where Claude 3.5 Sonnet was

System	PDF+VLM parsing	Code gen.	Code exec.	Metric vs. claim	Critic loop	Struct. report
MLAgentBench (2023)	✗	~	✓	✗	~	✗
PaperBench (2025)	✓	✓	✓	~	~	✗
PaperCoder (2025)	✓	✓	✗	✗	✗	~
AutoP2C (2025)	✓	✓	✓	~	~	✗
AutoReproduce (2025)	✓	✓	✓	✗	~	✗
Agent Laboratory (2025)	~	✓	✓	✗	✓	✓
AI Scientist v1/v2 (2024)	✗/~	✓	✓	✗	✓	✓
MLR-Copilot (2024)	✗	✓	✓	✗	~	✗
ReproBot (ours)	✓	✓	✓	✓	✓	✓

Table 1: Capability gap matrix across replication and research-generation systems. ✓ = full support, ~ = partial, ✗ = not supported. The critical missing column in every prior system is *metric vs. claim* paired with a Critic retry loop – exactly what ReproBot adds. The specific design choices borrowed from each system are detailed in Section 3.




Figure 1: ReproBot pipeline. The Orchestrator drives the Reader → Coder → Runner → Critic loop until results clear a confidence threshold or the retry budget is exhausted, then hands off to the Report Generator.

the clear leader among tested models (21.0% average replication score vs. GPT-4o’s 4.1% and o1’s 13.2%) on exactly this kind of paper-comprehension-and-agentic-coding task [2].

4 Evaluation Methodology

The four replication-focused systems reviewed in Section 2.3 each invented their own evaluation vocabulary. Table 2 maps these onto a shared taxonomy and proposes a combined rubric for ReproBot’s Critic: PaperBench’s three-tier taxonomy is adopted as the skeleton (it is the most validated, with an author-co-developed rubric and a judge-reliability meta-check reaching $F1 = 0.83$ [2]), filled with the cheaper, more mechanistic techniques the other systems proved effective, plus the one thing none of

them define: an explicit numeric tolerance gate.

Success will be reported along three axes rather than a single aggregate replication rate: **coverage** (how many target papers produced any executing script), **fidelity** (of those, how many reproduced the claimed metric within the stated tolerance, with the actual gap percentage reported per paper), and, as the headline result, an **ablation** of replication rate with the Critic-driven retry loop enabled versus disabled (single-shot Coder only) – isolating the specific contribution the literature review identifies as ReproBot’s gap-filling feature.

5 Project Timeline

The original proposal laid out four monthly milestones: Reader (Month 1), Coder + Runner + Or-

Tier	What it checks	Borrowed technique	Output
Code Development	Does the script structurally implement the method?	Component-coverage check by section (method / data-processing / eval)	partial-credit %
Execution	Does it run, cheaply verified before a full run?	Sampling-based dry-run (mini-batch check + early-exit pipeline pass)	binary + triage class
Result Match	Does the reproduced metric match the paper’s claim?	Continuous gap signal feeding an explicit numeric tolerance threshold	pass / retry / fail
Meta-layer	Is the Critic’s own judgment trustworthy?	Validate Critic feedback against a small hand-graded gold set	judge F1 / correlation

Table 2: Proposed combined evaluation rubric for ReproBot’s Critic agent.

chestrator skeleton (Month 2), Critic + refinement loop + report generation (Month 3), and evaluation + ablation + demo (Month 4). Detailed planning (Section 6) refined this into the staged breakdown in Table 3, with an explicit pilot checkpoint – validating the pipeline end to end on 3–5 papers – before committing to the full 20-paper benchmark.

6 Progress to Date

This section is the core contribution of the present report: a concrete account of what has been completed since the project proposal, and what remains open.

6.1 Completed: Literature Review

Nine closely related systems were surveyed in depth and organized into three internal documents that serve different purposes: a cross-paper comparison document with a landscape diagram, a chronological wave structure, per-paper architecture deep dives, and a capability gap matrix; a template-consistent per-paper summary document preserving exact reported numbers (success rates, replication scores, costs) for later citation; and a polished prose draft of the Introduction and Literature Review sections, ready to be copy-edited into the final project report (Section 2 of this document is drawn directly from it). A companion evaluation-metrics comparison document (Section 4) inventories every metric these four replication-focused papers report, checks judge/grader trustworthiness, and proposes the combined rubric ReproBot’s Critic will use.

6.2 Completed: Detailed Project Plan

Beyond the original proposal, a detailed, implementation-ready project plan was produced, covering: a feasibility assessment grounded in comparable systems’ actual reported numbers (Table 4 reproduces the key comparison); an architecture deep dive with the concrete shared-memory JSON schema threaded through every agent call; per-agent

specifications (responsibilities, tools, recommended model, literature-grounded risks and mitigations); the revised staged timeline in Section 5; a risk register; and a rough compute/cost budget.

6.3 Early Prototyping: Reader Agent

An early prototype of the **Reader** agent has been sketched out, but it is at a very preliminary stage. It uses **pdfplumber** to extract text from a paper PDF and the Anthropic API (Claude) to draft a mostly-structured method summary, claims, and hyperparameters from that text.

So far it has only been tried on a small handful of CIFAR-10 image-classification papers (ResNet, DenseNet, and similar canonical architectures) – deliberately chosen to match the project’s launch domain. This is a first sanity check of the extraction idea, not a validated component: no systematic accuracy evaluation has been run, the output schema is not yet reliably machine-parseable, and the vision-language component for figures, diagrams, and tables is not yet integrated (the prototype relies on **pdfplumber** text alone).

AutoP2C’s ablation is a direct reminder that the VLM piece matters: removing architecture-diagram parsing alone cost 21.9 points of absolute reproduced performance in their evaluation [4]. Hardening the schema and adding VLM-based figure parsing are therefore near-term priorities, not optional enhancements.

6.4 Not Yet Started

The Coder, Runner, and Critic agents, the Orchestrator’s shared-memory state machine, the Docker sandbox, and the report generator have not yet been implemented. No pilot-paper validation (the 3–5 paper checkpoint in Table 3) has been run yet, and consequently no reproduced metrics, executability rates, or ablation results exist at this stage. The project team has also grown from the two authors on the original proposal to four members (Aleyna Küçük, Engincan Varan, Gül Korkut, Göktuğ Mert Özdoğan), which has not yet been reflected in a re-

Phase	Focus
Month 1	PDF ingestion, Reader prompt/schema design, shared-memory object definition; Reader validated on 3–5 pilot papers.
Month 2	Coder generating scripts for the pilot set; Docker sandbox with sampling-based dry-run validation; Orchestrator skeleton (Reader → Coder → Runner) with a hard-coded single retry.
Month 3	Critic verdict logic (pass/retry/fail + three-way Execution/Code-Dev/Result-Match split); real retry loop with targeted feedback; Report Generator; single-shot vs. iterative-loop ablation on the pilot set.
Month 4	Expand from the pilot set toward the full 20-paper target as compute/time allows; Gradio demo; ablation write-up; final project report.

Table 3: Revised, staged project timeline.

System	Team / effort	What they achieved	Cost
PaperBench agent	BasicA- Dedicated OpenAI team	21.0% avg. replication score (Claude 3.5 Sonnet), 20 ICML papers	~\$400/paper + \$66/paper grading
PaperCoder	Research team, months	45.1% on Code-Dev (no execution); 28.5% on execution+result-match subset	Not reported
AutoReproduce	Research team	48.5% on Code-Dev with lineage; 77–95% execution rate on own benchmark	Not reported
Human ML PhDs	8 PhDs, 48h each	~41.4% best-of-3	N/A

Table 4: Reality check against comparable systems – no reviewed system reaches close to faithful replication, which frames the 20-paper benchmark as a demonstration of the retry-loop contribution rather than a claim to beat these numbers.

vised division of labor beyond the initial kickoff presentation.

5. Finalize the division of labor across the now four-person team against the agent specifications in Section 3.

7 Next Steps

The immediate priorities, in order, are:

1. **Harden the Reader’s extraction schema** so that method summary, claims (with explicit source location), hyperparameters, and a dedicated data-processing field are all reliably machine-parseable, and validate it by hand against 3–5 pilot papers as specified in the project plan.
2. **Add VLM-based figure and table parsing** to the Reader, given the concrete evidence in Section 2.3 that this materially affects reproduced performance.
3. **Define the shared-memory schema in code** and stand up a minimal Orchestrator skeleton that can pass the Reader’s output to a (initially stubbed) Coder.
4. **Begin the Coder agent**, targeting a single self-contained HuggingFace `Trainer` script per paper, and the Docker sandbox with the sampling-based dry-run check built first.

8 Conclusion

This report covers ReproBot’s first phase of work: a literature review of nine closely related agentic-ML and paper-replication systems, a detailed implementation-ready project plan, and an early prototype of the Reader agent. The literature review confirms a specific, well-evidenced gap – no surveyed system closes the loop from a reproduced metric to an explicit pass/fail verdict against the paper’s own claim to a targeted retry – that ReproBot’s Critic-driven design is intended to fill.

Implementation is still at an early stage. The Reader prototype has only been tried on a small sample of CIFAR-10 papers and is not yet schema-complete or multimodal, and the Coder, Runner, Critic, and Orchestrator remain to be built. The next phase focuses on maturing and validating the Reader against the pilot paper set before extending the pipeline to code generation and execution.

References

- [1] Huang, Q., Vora, J., Liang, P., & Leskovec, J. (2023). MLAGentBench: Evaluating Language Agents on Machine Learning Experimentation. *arXiv:2310.03302*.
- [2] Starace, G., Jaffe, O., Sherburn, D., Aung, J., et al. (2025). PaperBench: Evaluating AI’s Ability to Replicate AI Research. *arXiv:2504.01848*.
- [3] Seo, M., Baek, J., Lee, S., & Hwang, S. J. (2025). Paper2Code: Automating Code Generation from Scientific Papers in Machine Learning. *arXiv:2504.17192* (ICLR 2026).
- [4] Lin, Z., Shen, Y., Cai, Q., Sun, H., Zhou, J., & Xiao, M. (2025). AutoP2C: An LLM-Based Agent Framework for Code Repository Generation from Multimodal Content in Academic Papers. *arXiv:2504.20115*.
- [5] Zhao, X., Sang, Z., Li, Y., Shi, Q., et al. (2025). AutoReproduce: Automatic AI Experiment Reproduction with Paper Lineage. *arXiv:2505.20662*.
- [6] Schmidgall, S., Su, Y., et al. (2025). Agent Laboratory: Using LLM Agents as Research Assistants. *arXiv:2501.04227*.
- [7] Lu, C., Lu, C., Lange, R. T., Foerster, J., Clune, J., & Ha, D. (2024). The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery. *arXiv:2408.06292*.
- [8] Yamada, Y., Lange, R. T., Lu, C., et al. (2025). The AI Scientist-v2: Workshop-Level Automated Scientific Discovery via Agentic Tree Search. *arXiv:2504.08066*.
- [9] Li, R., Patel, T., Wang, Q., & Du, X. (2024). MLR-Copilot: Autonomous Machine Learning Research based on Large Language Models Agents. *arXiv:2408.14033*.